

OSM

G assignment 2

Tobias Hallundbæk Petersen (xtv657)

Ola Rønning (vdl761)

Nikolaj Høyer (ctl533)

February 24, 2014

Contents

1	Types and Functions for Userland Processes in Buenos	2
1.1	Define a data structure to represent a user process	2
1.2	Implement a library of helper functions	2
1.2.1	process_spawn	3
1.2.2	process_finish	4
1.2.3	process_join	5
1.2.4	process_init	6
2	System Calls for User-Process Control in Buenos	6
3	Testing	7

1 Types and Functions for Userland Processes in Buenos

1.1 Define a data structure to represent a user process

We start by defining the data structure needed for representing user processes. This is done in `syscall.h`. First we define a process state as consisting of a number of different states, then we define our process control block as using this process state, along with a process id (an integer), a parent id, a name, a return value and an integer representing the number of child processes.

```
#define CONFIG_MAX_NAME 128
typedef enum {
    PROC_FREE,
    PROC_RUNNING,
    PROC_READY,
    PROC_SLEEPING,
    PROC_ZOMBIE,
    PROC_NONREADY,
    PROC_DYING,
    PROC_NOTFREE
} process_state_t;

typedef struct {
    process_state_t state;
    process_id_t id;
    process_id_t parentid;
    char name[CONFIG_MAX_NAME];
    int retval;
    int children;
} process_control_block_t;
```

1.2 Implement a library of helper functions

We start by modifying the existing `process_start` function, making it take a process id as input, instead of an executable. We make sure that the thread entry `my_entry` gets the process id of the input, and also we modify the `file = ...` line to update the process table with the process name, as well as opening the file. This also requires a spinlock, to prevent corruption of the process table.

```
void process_start(process_id_t pid)
{
    ...
    my_entry->process_id = pid;
```

```

...
spinlock_acquire(&process_lock);
file = vfs_open((char *)process_table[pid].name);
spinlock_release(&process_lock);
...

```

We also need a helper function, which creates a new process id. Such a function is defined below. We look through the process table and checks for a free process. This is then returned. Also, a spinlock is required when accessing the process table. If this in some way or another fails, -1 is returned.

```

process_id_t process_new_id() {
    int pid;
    process_control_block_t process;
    spinlock_acquire(&process_lock);
    for (pid = 0; pid < CONFIG_MAX_PROCESSES; pid++) {
        process = process_table[pid];
        if (process.state == PROC_FREE){
            process.parentid = -1;
            process.children = 0;
            process.state = PROC_NOTFREE;
            spinlock_release(&process_lock);
            return pid;
        }
    }
    spinlock_release(&process_lock);
    return -1;
}

```

1.2.1 process_spawn

When spawning a process, we give an executable as input and return its process id. Thus, we need the helper function described above - we start by acquiring an available process id, then update our process table. Again, we apply a spinlock before making changes to the table. The parent process for the new process must be the current running process, hence `parentid` is set with the value of `process_get_current_process()`. We also use `stringcopy` to place the name of the executable in the table. The other table values should be clear from the context.

We then make a check - if `pid` is 0, then it is our init process running. If this is the case, we simply set the process as running and return. Otherwise, its some sub-process and we create a new thread for the child process and run it. Also, we increment the parents child-counter.

```

process_id_t process_spawn(const char *executable) {
    process_id_t pid;
    TID_t child_tid;
    pid = process_new_id();

    interrupt_status_t intr_status;
    intr_status = _interrupt_disable();
    spinlock_acquire(&process_lock);
    process_id_t parent_process = process_get_current_process();
    process_table[pid].state = PROC_READY;
    process_table[pid].parentid = parent_process;
    process_table[pid].id = pid;
    process_table[pid].children = 0;
    strcpy(process_table[pid].name, executable, CONFIG_MAX_NAME);
    if (pid == 0) {
        process_table[pid].state = PROC_RUNNING;
        spinlock_release(&process_lock);
        _interrupt_set_state(intr_status);
        return pid;
    }
    child_tid = thread_create((void (*)(uint32_t))process_start, (uint32_t) pid)↵
    ;
    if (child_tid < 0){
        process_table[pid].state = PROC_FREE;
        spinlock_release(&process_lock);
        _interrupt_set_state(intr_status);
        return -1;
    }
    process_table[parent_process].children += 1;
    process_table[pid].state = PROC_RUNNING;
    spinlock_release(&process_lock);
    _interrupt_set_state(intr_status);
    thread_run(child_tid);
    return pid;
}

```

1.2.2 process_finish

The process-finish function gets the current process, sets it as a ZOMBIE and sets its return value. It then wakes up all threads for the current process. This requires both interrupts to be disabled, since we need to use to sleep queue, and a spin lock, since we need to access the process table.

We then destroy the page table of the thread and finishes the thread.

```

/* Stop the process and the thread it runs in. Sets the return value as well ↵
 */
void process_finish(int retval) {
    thread_table_t *thr;
    interrupt_status_t intr_status;

```

```

    intr_status = _interrupt_disable();
    spinlock_acquire(&process_lock);
    process_id_t current_process = process_get_current_process();
    process_table[current_process].state = PROC_ZOMBIE;
    process_table[current_process].retval = retval;
    sleepq_wake_all(&(process_table[current_process]));
    spinlock_release(&process_lock);
    _interrupt_set_state(intr_status);

    thr = thread_get_current_thread_entry();
    vm_destroy_pagetable(thr->pagetable);
    thr->pagetable = NULL;
    thread_finish();
}

```

1.2.3 process_join

The process-join function disables interrupts whilst waiting for the child process to finish, this function can only be called by parent process. The child process is added to the sleep queue and the threads are switched, which is continued till the child process has finished. When the child process has finished its return value is stored and the child process is freed. The return value of the process is returned by the join function.

```

int process_join(process_id_t pid) {
    interrupt_status_t intr_status;
    int retval;

    intr_status = _interrupt_disable();
    spinlock_acquire(&process_lock);

    if (process_get_current_process() != process_table[pid].parentid) return -1;
    while (process_table[pid].state != PROC_DYING) {
        sleepq_add(&(process_table[pid]));
        spinlock_release(&process_lock);
        thread_switch();
        spinlock_acquire(&process_lock);
    }

    retval = process_table[pid].retval;
    spinlock_release(&process_lock);
    _interrupt_set_state(intr_status);

    return retval;
}

```

1.2.4 process_init

Process init simply runs through the process table, setting its state to **FREE** and initialising all other values its values to defaults as well. It also resets the spinlock, to make sure it has its default state.

```
void process_init() {
    int i;
    spinlock_reset(&process_lock);
    for (i=0; i<CONFIG_MAX_PROCESSES; i++) {
        process_table[i].state      = PROC_FREE;
        process_table[i].id        = -1;
        process_table[i].parentid  = -1;
        process_table[i].name[0]   = '\0';
        process_table[i].retval    = -1;
        process_table[i].children  = 0;
    }
}
```

2 System Calls for User-Process Control in Buenos

We now modify `syscall.c` to use the functions implemented in `process.c`. We simply add functions we can call when the appropriate syscalls are invoked (see below).

```
int syscall_exec(char const* filename){
    int pid = process_spawn(filename);
    return pid;
}
void syscall_exit(int retval){
    process_finish(retval);
}
int syscall_join(int pid){
    return process_join(pid);
}
```

We now add cases to the syscalls in `syscall.c`, that loads and writes to registers, calling the functions defined above.

```
switch(user_context->cpu_regs[MIPS_REGISTER_A0])
...
case SYSCALL_EXEC:
    V0 = syscall_exec((char *) A1);
    break;
case SYSCALL_EXIT:
```

```

        syscall_exit(A1);
        break;
case SYSCALL_JOIN:
    V0 = syscall_join(A1);
    break;
...

```

3 Testing

Apart from the included `exec` test, we have made another test-file, `multiexec`, which tests execution of multiple programs simultaneously, in this case two instances of the hello world program, `hw`. The test program is shown below:

```

/*
 * Userland exec test
 */

#include "tests/lib.h"

static const char prog0[] = "[buenos]tests/hw"; /* The program to start. */

int main(void)
{
    uint32_t child0, child1;
    int ret;
    printf("Starting program %s\n", prog0);
    child0 = syscall_exec(prog0);
    printf("Starting program %s\n", prog0);
    child1 = syscall_exec(prog0);
    printf("Now joining child %d\n", child0);
    ret = (char)syscall_join(child0);
    printf("Now joining child %d\n", child1);
    ret = (char)syscall_join(child1);
    printf("Child joined with status: %d\n", ret);
    syscall_halt();
    return 0;
}

```