

OSM

G assignment 2

Tobias Hallundbæk Petersen (xtv657)

Ola Rønning (vdl761)

Nikolaj Høyer (ctl533)

February 24, 2014

Contents

1	Types and Functions for Userland Processes in Buenos	2
1.1	Define a data structure to represent a user process	2
1.2	Implement a library of helper functions	2
1.2.1	<code>process_spawn</code>	3
1.2.2	<code>process_finish</code>	4
1.2.3	<code>process_join</code>	4
1.2.4	<code>process_init</code>	5
2	System Calls for User-Process Control in Buenos	5

1 Types and Functions for Userland Processes in Buenos

1.1 Define a data structure to represent a user process

We start by defining the data structure needed for representing user processes. This is done in `syscall.h`. First we define a process state as consisting of a number of different states, then we define our process control block as using this process state, along with a process id (an integer), a parent id, a name, a return value and an integer representing the number of child processes.

```
#define CONFIG_MAX_NAME 128
typedef enum {
    PROC_FREE,
    PROC_RUNNING,
    PROC_READY,
    PROC_SLEEPING,
    PROC_ZOMBIE,
    PROC_NONREADY,
    PROC_DYING,
    PROC_NOTFREE
} process_state_t;

typedef struct {
    process_state_t state;
    process_id_t id;
    process_id_t parentid;
    char name[CONFIG_MAX_NAME];
    int retval;
    int children;
} process_control_block_t;
```

1.2 Implement a library of helper functions

```
void process_start(process_id_t pid)
{
    thread_table_t *my_entry;
    pagetable_t *pagetable;
    uint32_t phys_page;
    context_t user_context;
    uint32_t stack_bottom;
    elf_info_t elf;
    openfile_t file;

    int i;
```

```

interrupt_status_t intr_status;

my_entry = thread_get_current_thread_entry();
/* If the pagetable of this thread is not NULL, we are trying to
   run a userland process for a second time in the same thread.
   This is not possible. */
KERNEL_ASSERT(my_entry->pagetable == NULL);

my_entry->process_id = pid;
pagetable = vm_create_pagetable(thread_get_current_thread());
KERNEL_ASSERT(pagetable != NULL);

intr_status = _interrupt_disable();
my_entry->pagetable = pagetable;
_interrupt_set_state(intr_status);
spinlock_acquire(&process_lock);
file = vfs_open((char *)process_table[pid].name);
spinlock_release(&process_lock);

```

```

process_id_t process_new_id() {
    int pid;
    process_control_block_t process;
    spinlock_acquire(&process_lock);
    for (pid = 0; pid < CONFIG_MAX_PROCESSES; pid++) {
        process = process_table[pid];
        if (process.state == PROC_FREE || (process.state == PROC_DYING &&
            (process.parentid < 0 || (process_table[process.parentid].state == PROC_DYING &&
                process.children == 0)))){
            process.parentid = -1;
            process.children = 0;
            process.state = PROC_NOTFREE;
            spinlock_release(&process_lock);
            return pid;
        }
    }
    spinlock_release(&process_lock);
    return -1;
}

```

1.2.1 process_spawn

```

process_id_t process_spawn(const char *executable) {
    process_id_t pid;
    TID_t child_tid;
    pid = process_new_id();

    spinlock_acquire(&process_lock);

```

```

process_id_t parent_process = process_get_current_process();
process_table[pid].state = PROC_READY;
process_table[pid].parentid = parent_process;
process_table[pid].id = pid;
process_table[pid].children = 0;
strcpy(process_table[pid].name, executable, CONFIG_MAX_NAME);
if (pid == 0) {
    process_table[pid].state = PROC_RUNNING;
    spinlock_release(&process_lock);
    return pid;
}
child_tid = thread_create((void (*)(uint32_t))process_start, (uint32_t) pid) ←
;
if (child_tid < 0){
    process_table[pid].state = PROC_FREE;
    return -1;
}
process_table[parent_process].children += 1;
process_table[pid].state = PROC_RUNNING;
spinlock_release(&process_lock);
thread_run(child_tid);
return pid;
}

```

1.2.2 process_finish

```

/* Stop the process and the thread it runs in. Sets the return value as well ←
*/
void process_finish(int retval) {
    thread_table_t *thr;
    interrupt_status_t intr_status;

    intr_status = _interrupt_disable();
    spinlock_acquire(&process_lock);
    process_id_t current_process = process_get_current_process();
    process_table[current_process].state = PROC_DYING;
    process_table[current_process].retval = retval;
    sleepq_wake_all(&(process_table[current_process]));
    spinlock_release(&process_lock);
    _interrupt_set_state(intr_status);

    thr = thread_get_current_thread_entry();
    vm_destroy_pagetable(thr->pagetable);
    thr->pagetable = NULL;
    thread_finish();
}

```

1.2.3 process_join

```

int process_join(process_id_t pid) {
    interrupt_status_t intr_status;
    int retval;

    intr_status = _interrupt_disable();
    spinlock_acquire(&process_lock);

    if (process_get_current_process() != process_table[pid].parentid) return -1;
    while (process_table[pid].state != PROC_DYING) {
        sleepq_add(&(process_table[pid]));
        spinlock_release(&process_lock);
        thread_switch();
        spinlock_acquire(&process_lock);
    }

    retval = process_table[pid].retval;
    spinlock_release(&process_lock);
    _interrupt_set_state(intr_status);

    return retval;
}

```

1.2.4 process_init

Process init simply runs through the process table, setting its state to **FREE** and initialising all other values its values to defaults as well. It also resets the spinlock, to make sure it has its default state.

```

void process_init() {
    int i;
    spinlock_reset(&process_lock);
    for (i=0; i<CONFIG_MAX_PROCESSES; i++) {
        process_table[i].state      = PROC_FREE;
        process_table[i].id         = -1;
        process_table[i].parentid   = -1;
        process_table[i].name[0]    = '\0';
        process_table[i].retval     = -1;
        process_table[i].children   = 0;
    }
}

```

2 System Calls for User-Process Control in Buenos

We now modify `syscall.c` to use the functions implemented in `process.c`. We simply add functions we can call when the appropriate syscalls are invoked (see below).

```

int syscall_exec(char const* filename){
    int pid = process_spawn(filename);
    return pid;
}
void syscall_exit(int retval){
    process_finish(retval);
}
int syscall_join(int pid){
    return process_join(pid);
}

```

We now add cases to the syscalls in `syscall.c`, that loads and writes to registers, calling the functions defined above.

```

switch(user_context->cpu_regs[MIPS_REGISTER_A0])
...
case SYSCALL_EXEC:
    V0 = syscall_exec((char *) A1);
    break;
case SYSCALL_EXIT:
    syscall_exit(A1);
    break;
case SYSCALL_JOIN:
    V0 = syscall_join(A1);
    break;
...

```